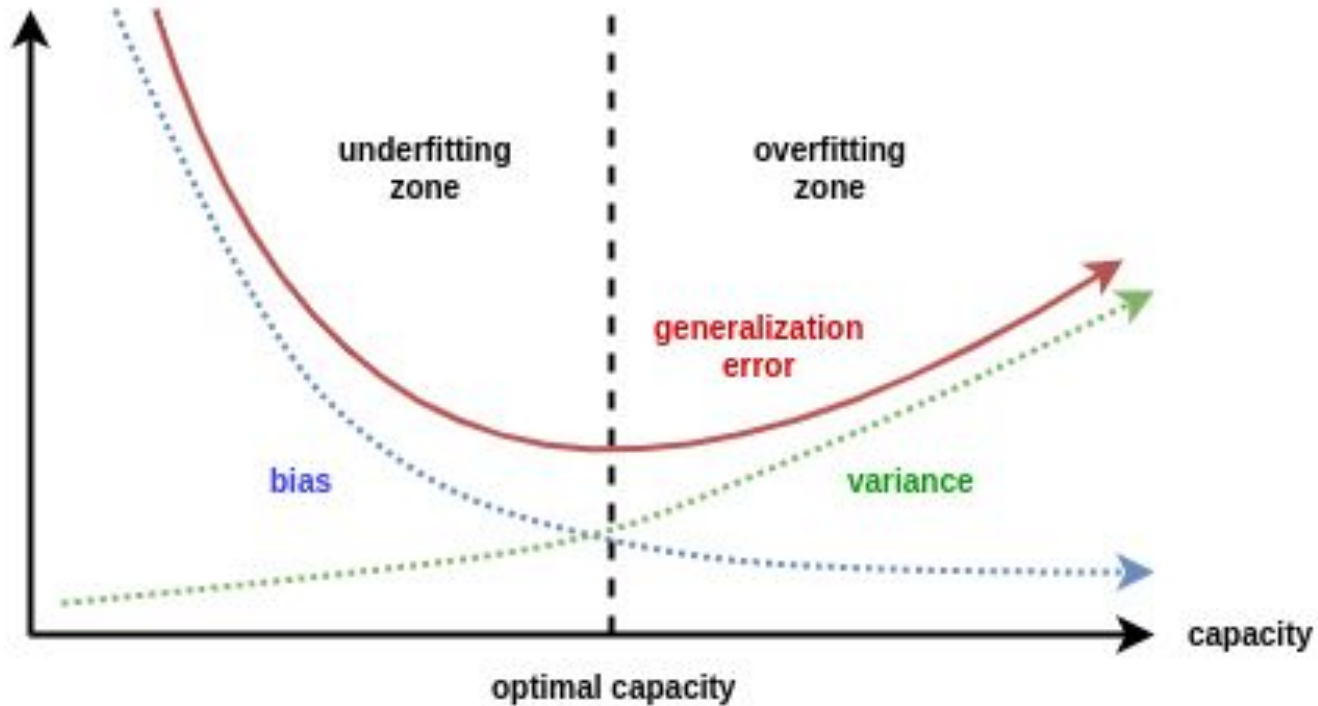# Lecture 8: Bias-Variance tradeoff; More Ensembling

Outline

1. Bias-Variance Tradeoff
2. Validation Strategies
3. Blending
4. Stacking

# Bias-Variance tradeoff

Bias-variance tradeoff

underfitting zone — overfitting zone

generalization error

bias

variance

optimal capacity

capacity

Low Variance High Variance

Low Bias

High Bias

Bias-variance decomposition derivation

Bias-variance decomposition

The dataset $X = (x_i, y_i)_{i=1}^{\ell}$ with $y_i \in \mathbb{R}$

for regression problem.

Denote loss function $L(y, a) = \big(y - a(x)\big)^2$.

The empirical risk takes form:

$$R(a) = \mathbb{E}_{x,y}\left[\big(y - a(x)\big)^2\right] = \int_{\mathbb{X}} \int_{\mathbb{Y}} p(x, y)\big(y - a(x)\big)^2 dx dy.$$

Bias-variance decomposition

Let's show that

$$a_*(x) = \mathbb{E}[y \,|\, x] = \int_{\mathbb{Y}} yp(y \,|\, x)dy = \arg\min_a R(a).$$

$$L(y, a(x)) = (y - a(x))^2 = (y - \mathbb{E}(y \,|\, x) + \mathbb{E}(y \,|\, x) - a(x))^2 =$$
$$= (y - \mathbb{E}(y \,|\, x))^2 + 2(y - \mathbb{E}(y \,|\, x))(\mathbb{E}(y \,|\, x) - a(x)) + (\mathbb{E}(y \,|\, x) - a(x))^2.$$

Returning to the risk estimation:

$$R(a) = \mathbb{E}_{x,y} L(y, a(x)) =$$
$$= \mathbb{E}_{x,y}(y - \mathbb{E}(y \,|\, x))^2 + \mathbb{E}_{x,y}(\mathbb{E}(y \,|\, x) - a(x))^2 +$$
$$+ 2\mathbb{E}_{x,y}(y - \mathbb{E}(y \,|\, x))(\mathbb{E}(y \,|\, x) - a(x)).$$

$$R(a) = \mathbb{E}_{x,y} L(y, a(x)) =$$

Focus on the last term:

$$= \mathbb{E}_{x,y}(y - \mathbb{E}(y \,|\, x))^2 + \mathbb{E}_{x,y}(\mathbb{E}(y \,|\, x) - a(x))^2 +$$

$$+ 2\mathbb{E}_{x,y}(y - \mathbb{E}(y \,|\, x))(\mathbb{E}(y \,|\, x) - a(x)).$$

Does not depend on y

$$\mathbb{E}_x \mathbb{E}_y \left[(y - \mathbb{E}(y \,|\, x)) \boxed{(\mathbb{E}(y \,|\, x) - a(x))} \,|\, x \right] =$$

$$= \mathbb{E}_x \left((\mathbb{E}(y \,|\, x) - a(x)) \mathbb{E}_y \left[(y - \mathbb{E}(y \,|\, x)) \,|\, x \right] \right) =$$

$$= \mathbb{E}_x \left((\mathbb{E}(y \,|\, x) - a(x))(\mathbb{E}(y \,|\, x) - \mathbb{E}(y \,|\, x)) \right) =$$
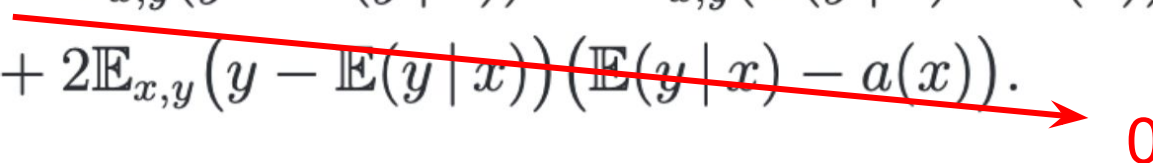
$$= 0$$

9

$$R(a) = \mathbb{E}_{x,y} L(y, a(x)) =$$

Focus on the last term:
$$= \mathbb{E}_{x,y}(y - \mathbb{E}(y \mid x))^2 + \mathbb{E}_{x,y}(\mathbb{E}(y \mid x) - a(x))^2 +$$
$$+ 2\mathbb{E}_{x,y}(y - \mathbb{E}(y \mid x))(\mathbb{E}(y \mid x) - a(x)).$$

0

$$\mathbb{E}_x \mathbb{E}_y \Big[(y - \mathbb{E}(y \mid x))(\mathbb{E}(y \mid x) - a(x)) \mid x \Big] =$$
$$= \mathbb{E}_x \Big((\mathbb{E}(y \mid x) - a(x)) \mathbb{E}_y \Big[(y - \mathbb{E}(y \mid x)) \mid x \Big] \Big) =$$
$$= \mathbb{E}_x \Big((\mathbb{E}(y \mid x) - a(x))(\mathbb{E}(y \mid x) - \mathbb{E}(y \mid x)) \Big) =$$
$$= 0$$

So the risk takes form:

$$R(a) = \boxed{\mathbb{E}_{x,y}(y - \mathbb{E}(y \mid x))^2} + \mathbb{E}_{x,y}(\mathbb{E}(y \mid x) - a(x))^2.$$

Does not depend on a(x)

The minimum is reached when $a(x) = \mathbb{E}(y \mid x).$

So the optimal regression model with square loss is

$$a_*(x) = \mathbb{E}(y \mid x) = \int_{\mathbb{Y}} y p(y \mid x) dy.$$

11

Denote $\mu : (\mathbb{X} \times \mathbb{Y})^\ell \to \mathcal{A}$, where $\mathcal{A}$ is some family of algorithms.

Denote $\mu : (\mathbb{X} \times \mathbb{Y})^\ell \to \mathcal{A}$, where $\mathcal{A}$ is some family of algorithms.

So $L(\mu) = \mathbb{E}_X \left[\mathbb{E}_{x,y} \left[(y - \mu(X)(x))^2 \right] \right]$, where X dataset.

In further slides (x) is omitted!

Denote $\mu : (\mathbb{X} \times \mathbb{Y})^\ell \to \mathcal{A}$, where $\mathcal{A}$ is some family of algorithms.

So $L(\mu) = \mathbb{E}_X\left[\mathbb{E}_{x,y}\left[(y - \mu(X)(x))^2\right]\right]$, where X dataset.

In further slides (x) is omitted!

If X is fixed, then

$$\mathbb{E}_{x,y}\left[(y - \mu(X))^2\right] = \mathbb{E}_{x,y}\left[(y - \mathbb{E}[y \mid x])^2\right] + \mathbb{E}_{x,y}\left[(\mathbb{E}[y \mid x] - \mu(X))^2\right].$$

Denote $\mu : (\mathbb{X} \times \mathbb{Y})^\ell \to \mathcal{A}$, where $\mathcal{A}$ is some family of algorithms.

So $L(\mu) = \mathbb{E}_X \left[\mathbb{E}_{x,y} \left[\left(y - \mu(X)(x) \right)^2 \right] \right]$, where X dataset.

In further slides (x) is omitted!

If X is fixed, then

$$\mathbb{E}_{x,y} \left[\left(y - \mu(X) \right)^2 \right] = \mathbb{E}_{x,y} \left[\left(y - \mathbb{E}[y \mid x] \right)^2 \right] + \mathbb{E}_{x,y} \left[\left(\mathbb{E}[y \mid x] - \mu(X) \right)^2 \right].$$

Let's combine the latter equations:

Denote $\mu : (\mathbb{X} \times \mathbb{Y})^\ell \to \mathcal{A}$, where $\mathcal{A}$ is some family of algorithms.

So $L(\mu) = \mathbb{E}_X \left[\mathbb{E}_{x,y} \left[(y - \mu(X)(x))^2 \right] \right]$, where X dataset.

In further slides (x) is omitted!

If X is fixed, then

$$\mathbb{E}_{x,y} \left[(y - \mu(X))^2 \right] = \mathbb{E}_{x,y} \left[(y - \mathbb{E}[y \mid x])^2 \right] + \mathbb{E}_{x,y} \left[(\mathbb{E}[y \mid x] - \mu(X))^2 \right].$$

Let's combine the latter equations:

$$L(\mu) = \mathbb{E}_X \left[\underbrace{\mathbb{E}_{x,y} \left[(y - \mathbb{E}[y \mid x])^2 \right]}_{\text{Does not depend on X}} + \mathbb{E}_{x,y} \left[(\mathbb{E}[y \mid x] - \mu(X))^2 \right] \right]$$

Does not depend on X

$$L(\mu) = \underbrace{\mathbb{E}_{x,y}\left[\left(y - \mathbb{E}[y \mid x]\right)^2\right]}_{\text{noise}} +$$

$$+ \underbrace{\mathbb{E}_x\left[\left(\mathbb{E}_X\left[\mu(X)\right] - \mathbb{E}[y \mid x]\right)^2\right]}_{\text{bias}} + \underbrace{\mathbb{E}_x\left[\mathbb{E}_X\left[\left(\mu(X) - \mathbb{E}_X\left[\mu(X)\right]\right)^2\right]\right]}_{\text{variance}}.$$

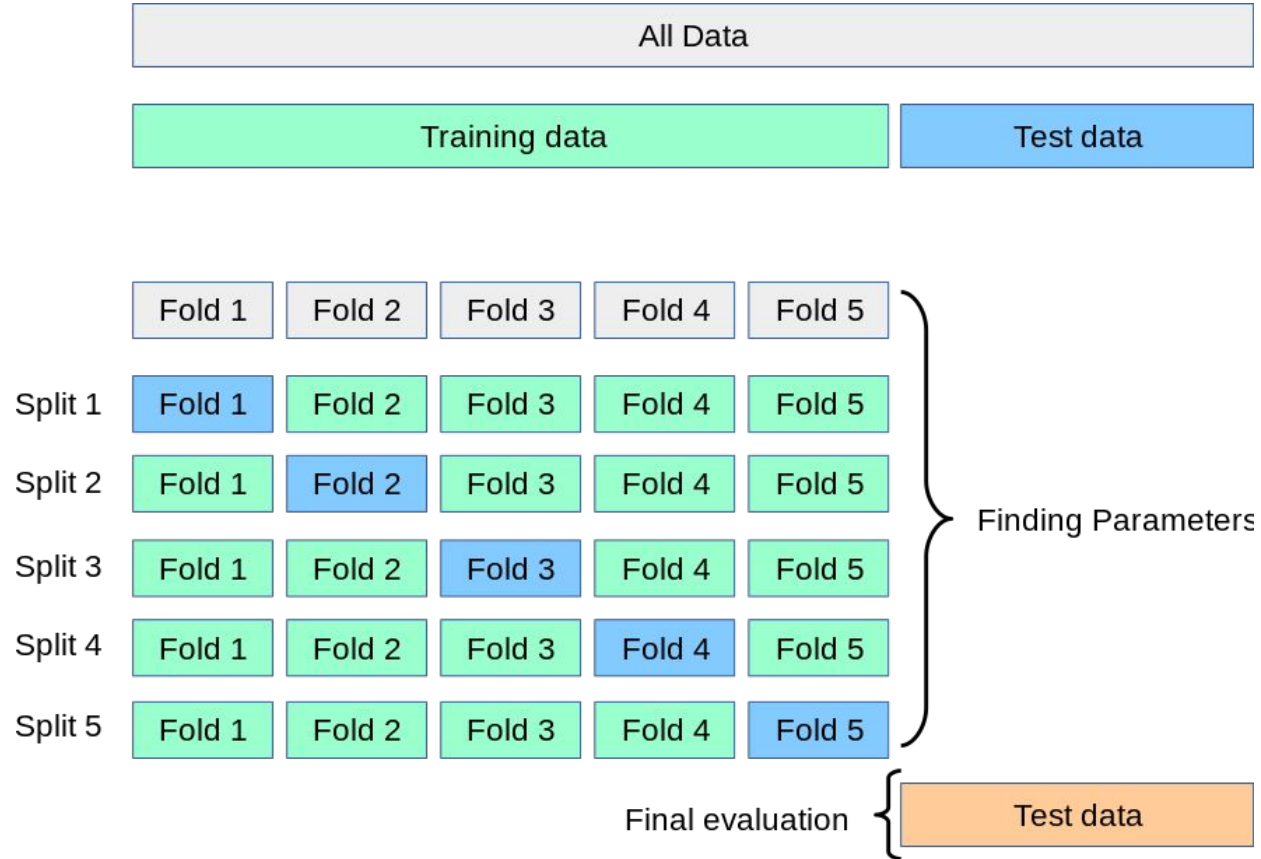Low Variance High Variance

Low Bias

High Bias

Bias-variance tradeoff

$$L(\mu) = \underbrace{\mathbb{E}_{x,y}\left[\left(y - \mathbb{E}[y \mid x]\right)^2\right]}_{\text{noise}} +$$

$$+ \underbrace{\mathbb{E}_x\left[\left(\mathbb{E}_X\left[\mu(X)\right] - \mathbb{E}[y \mid x]\right)^2\right]}_{\text{bias}} + \underbrace{\mathbb{E}_x\left[\mathbb{E}_X\left[\left(\mu(X) - \mathbb{E}_X\left[\mu(X)\right]\right)^2\right]\right]}_{\text{variance}}.$$

This exact form of bias-variance decomposition is correct for square loss in regression.

However, it is much more general. See extra materials for more exotic cases.

# Validation Strategies

Validation strategies

Validation strategies

Special case: Leave One Out (LOO) - good for tiny datasets

Validation strategies

Validation strategies

Special case: time series

Never use `train_test_split` in this case!

Stacking and blending

Blending

How to build an ensemble from *different* models?
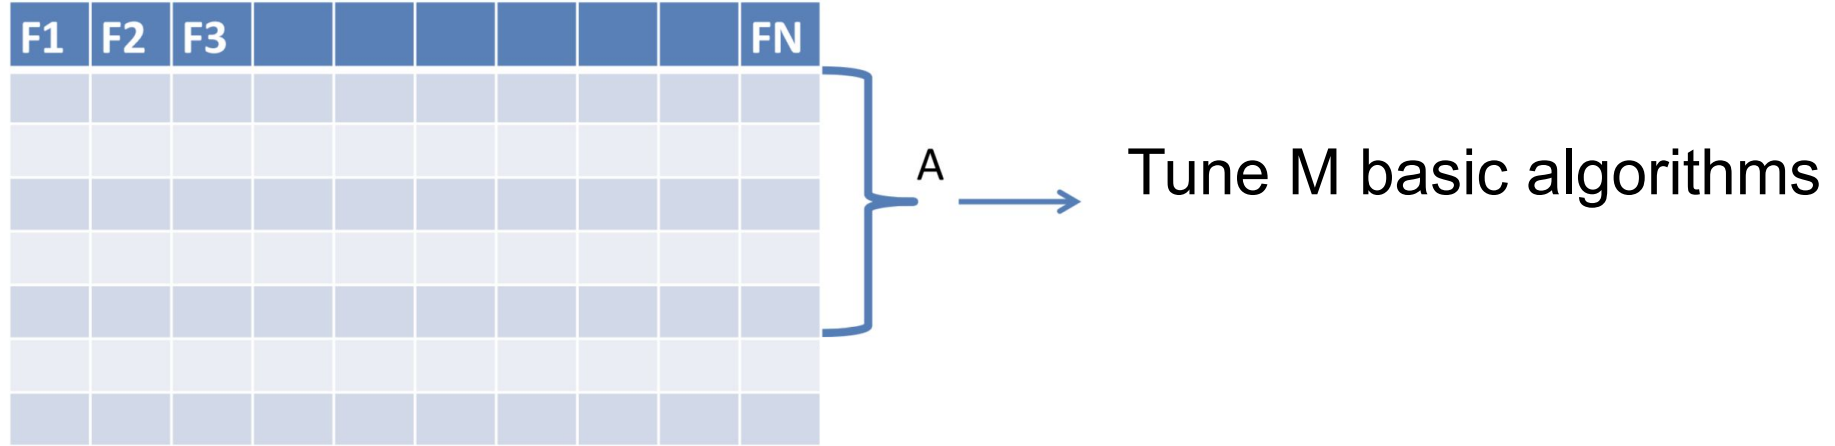
A → Tune M basic algorithms

Blending

How to build an ensemble from *different* models?

Tune M basic algorithms

Tune new model on the outputs of the basic algorithms

Blending

How to build an ensemble from *different* models?

Tune M basic algorithms
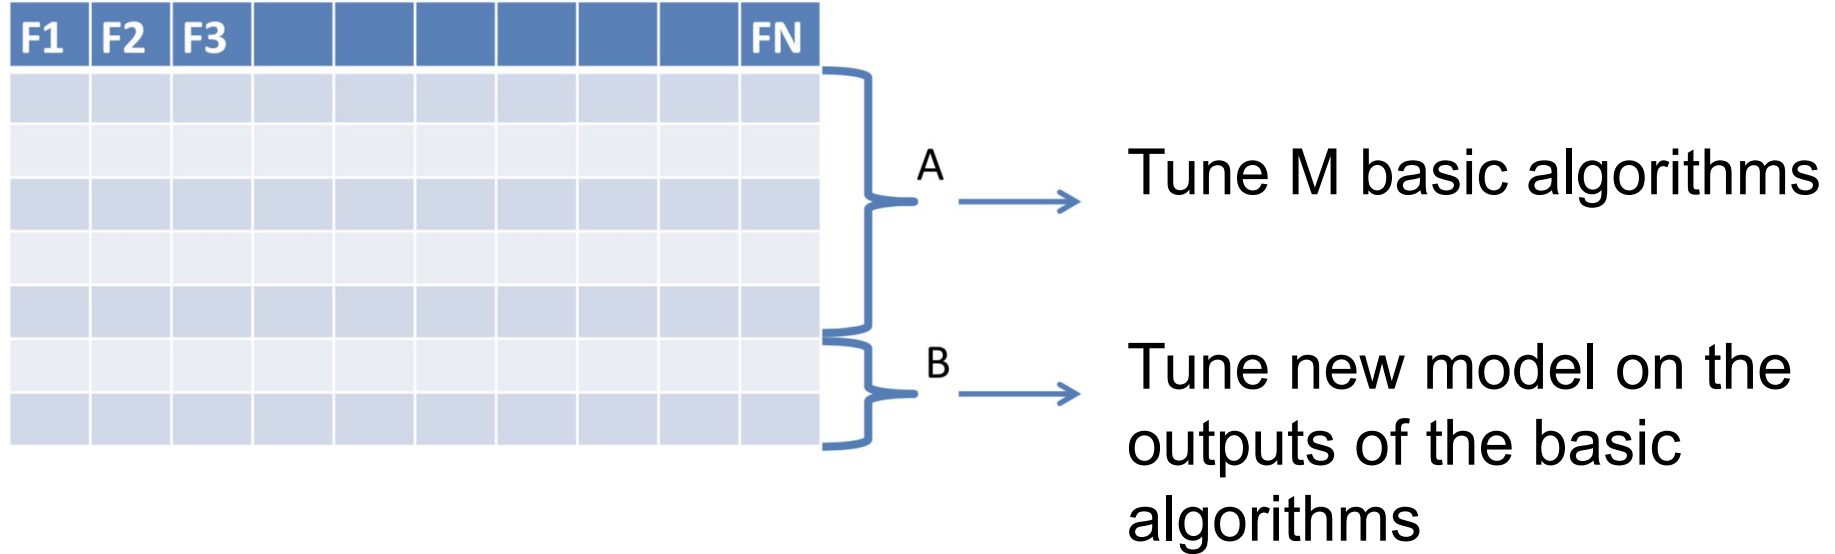
$$\hat{f}(x) = \sum_{i=1}^{M} \rho_i f_i(x) \qquad \sum_{i=1}^{M} \rho_i = 1, \quad \rho_i \in [0;1] \ \forall i$$

Blending

Just combine several *strong/complex* models.

$$\hat{f}(x) = \sum_{i=1}^{M} \rho_i f_i(x), \qquad \sum_{i=1}^{M} \rho_i = 1, \quad \rho_i \in [0; 1] \ \ \forall i$$

- Pros:
  - Simple and intuitive ensembling method.
  - Average several blendings to achieve better results.
- Cons:
  - Linear composition is not always enough.
  - Need to split the data. How to fix it?

Stacking

1. Split data into folds

Fit using folds to get meta-features on train

2. Tune models on different groups of folds, predict on left out

3. Tune the new model on the "meta"-features

Fit using all data meta-features on test

Picture source: https://dyakonov.org/2017/03/10/стекинг-stacking-и-блендинг-blending/

Stacking

- Train base algorithm(s) on different groups of folds leaving one fold out.

- Predict the meta-features on the left-out fold and test data.

- Train the meta-algorithm on the meta-features representation of the train data.

- Use it on the meta-features representation of the test data.

Stacking

- Pros:
  - Powerful ensembling method, if you know how to use it
  - Quite popular in ML-competitions
  - One might perform stacking on the meta-features dataset as well
- Cons:
  - Meta-features on each fold are actually predicted by different models
    - However, regularization usually helps
  - Hard to explain your model behaviour

Recap: ensembling methods

1. Bagging.
2. Random subspace method (RSM).
3. Bagging + RSM + Decision trees = Random Forest.
4. Blending.
5. Stacking.

Q & A

Appendix
(continue from slide 13)

$$L(\mu) = \mathbb{E}_X \left[\underbrace{\mathbb{E}_{x,y} \left[\left(y - \mathbb{E}[y \mid x] \right)^2 \right]}_{\text{Does not depend on X}} + \mathbb{E}_{x,y} \left[\left(\mathbb{E}[y \mid x] - \mu(X) \right)^2 \right] \right] =$$

Does not depend on X

$$L(\mu) = \mathbb{E}_X \Bigg[\underbrace{\mathbb{E}_{x,y}\Big[\big(y - \mathbb{E}[y \mid x]\big)^2\Big]}_{\text{Does not depend on X}} + \mathbb{E}_{x,y}\Big[\big(\mathbb{E}[y \mid x] - \mu(X)\big)^2\Big]\Bigg] =$$

$$= \mathbb{E}_{x,y}\Big[\big(y - \mathbb{E}[y \mid x]\big)^2\Big] + \mathbb{E}_{x,y}\Big[\mathbb{E}_X\Big[\big(\mathbb{E}[y \mid x] - \mu(X)\big)^2\Big]\Big].$$

$$L(\mu) = \mathbb{E}_X \left[\mathbb{E}_{x,y} \left[\left(y - \mathbb{E}[y \mid x] \right)^2 \right] + \mathbb{E}_{x,y} \left[\left(\mathbb{E}[y \mid x] - \mu(X) \right)^2 \right] \right] =$$

$$= \mathbb{E}_{x,y} \left[\left(y - \mathbb{E}[y \mid x] \right)^2 \right] + \boxed{\mathbb{E}_{x,y} \left[\mathbb{E}_X \left[\left(\mathbb{E}[y \mid x] - \mu(X) \right)^2 \right] \right]}.$$

Focus on the second term:

$$L(\mu) = \mathbb{E}_X \left[\mathbb{E}_{x,y} \left[(y - \mathbb{E}[y \mid x])^2 \right] + \mathbb{E}_{x,y} \left[(\mathbb{E}[y \mid x] - \mu(X))^2 \right] \right] =$$

$$= \mathbb{E}_{x,y} \left[(y - \mathbb{E}[y \mid x])^2 \right] + \boxed{\mathbb{E}_{x,y} \left[\mathbb{E}_X \left[(\mathbb{E}[y \mid x] - \mu(X))^2 \right] \right]}.$$

Focus on the second term:

$$\mathbb{E}_{x,y} \left[\mathbb{E}_X \left[(\mathbb{E}[y \mid x] - \mu(X))^2 \right] \right] =$$

$$L(\mu) = \mathbb{E}_X\left[\mathbb{E}_{x,y}\left[(y - \mathbb{E}[y \mid x])^2\right] + \mathbb{E}_{x,y}\left[(\mathbb{E}[y \mid x] - \mu(X))^2\right]\right] =$$

$$= \mathbb{E}_{x,y}\left[(y - \mathbb{E}[y \mid x])^2\right] + \mathbb{E}_{x,y}\left[\mathbb{E}_X\left[(\mathbb{E}[y \mid x] - \mu(X))^2\right]\right].$$

Focus on the second term:

$$\mathbb{E}_{x,y}\left[\mathbb{E}_X\left[(\mathbb{E}[y \mid x] - \mu(X))^2\right]\right] =$$

$$= \mathbb{E}_{x,y}\left[\mathbb{E}_X\left[(\mathbb{E}[y \mid x] - \mathbb{E}_X\left[\mu(X)\right] + \mathbb{E}_X\left[\mu(X)\right] - \mu(X))^2\right]\right]$$

$$\mathbb{E}_{x,y}\left[\mathbb{E}_X\left[\left(\mathbb{E}[y\mid x] - \mu(X)\right)^2\right]\right] =$$

$$= \mathbb{E}_{x,y}\left[\mathbb{E}_X\left[\left(\mathbb{E}[y\mid x] - \mathbb{E}_X\left[\mu(X)\right] + \mathbb{E}_X\left[\mu(X)\right] - \mu(X)\right)^2\right]\right] =$$

$$\mathbb{E}_{x,y}\left[\mathbb{E}_X\left[\left(\mathbb{E}[y\mid x]-\mu(X)\right)^2\right]\right]=$$

$$=\mathbb{E}_{x,y}\left[\mathbb{E}_X\left[\left(\mathbb{E}[y\mid x]-\mathbb{E}_X\left[\mu(X)\right]+\mathbb{E}_X\left[\mu(X)\right]-\mu(X)\right)^2\right]\right]=$$

$$=\mathbb{E}_{x,y}\left[\mathbb{E}_X\left[\underbrace{\left(\mathbb{E}[y\mid x]-\mathbb{E}_X\left[\mu(X)\right]\right)^2}\right]\right]+\mathbb{E}_{x,y}\left[\mathbb{E}_X\left[\left(\mathbb{E}_X\left[\mu(X)\right]-\mu(X)\right)^2\right]\right]+$$

$$+2\mathbb{E}_{x,y}\left[\mathbb{E}_X\left[\left(\mathbb{E}[y\mid x]-\mathbb{E}_X\left[\mu(X)\right]\right)\left(\mathbb{E}_X\left[\mu(X)\right]-\mu(X)\right)\right]\right].$$

$$\mathbb{E}_{x,y}\left[\mathbb{E}_X\left[\left(\mathbb{E}[y\mid x]-\mu(X)\right)^2\right]\right]=$$

$$=\mathbb{E}_{x,y}\left[\mathbb{E}_X\left[\left(\mathbb{E}[y\mid x]-\mathbb{E}_X\left[\mu(X)\right]+\mathbb{E}_X\left[\mu(X)\right]-\mu(X)\right)^2\right]\right]=$$

$$=\mathbb{E}_{x,y}\left[\mathbb{E}_X\left[\underbrace{\left(\mathbb{E}[y\mid x]-\mathbb{E}_X\left[\mu(X)\right]\right)^2}\right]\right]+\mathbb{E}_{x,y}\left[\mathbb{E}_X\left[\left(\mathbb{E}_X\left[\mu(X)\right]-\mu(X)\right)^2\right]\right]+$$

Does not depend on X

$$+2\mathbb{E}_{x,y}\left[\mathbb{E}_X\left[\left(\mathbb{E}[y\mid x]-\mathbb{E}_X\left[\mu(X)\right]\right)\left(\mathbb{E}_X\left[\mu(X)\right]-\mu(X)\right)\right]\right].$$

$$\mathbb{E}_{x,y}\left[\mathbb{E}_X\left[\left(\mathbb{E}[y\mid x]-\mu(X)\right)^2\right]\right]=$$

$$=\mathbb{E}_{x,y}\left[\mathbb{E}_X\left[\left(\mathbb{E}[y\mid x]-\mathbb{E}_X\left[\mu(X)\right]+\mathbb{E}_X\left[\mu(X)\right]-\mu(X)\right)^2\right]\right]=$$

$$=\mathbb{E}_{x,y}\left[\mathbb{E}_X\left[\underbrace{\left(\mathbb{E}[y\mid x]-\mathbb{E}_X\left[\mu(X)\right]\right)^2}\right]\right]+\mathbb{E}_{x,y}\left[\mathbb{E}_X\left[\left(\mathbb{E}_X\left[\mu(X)\right]-\mu(X)\right)^2\right]\right]+$$

Does not depend on X

$$+2\mathbb{E}_{x,y}\left[\mathbb{E}_X\left[\left(\mathbb{E}[y\mid x]-\mathbb{E}_X\left[\mu(X)\right]\right)\left(\mathbb{E}_X\left[\mu(X)\right]-\mu(X)\right)\right]\right].$$

Just a bit further, we are almost there

$$\mathbb{E}_{x,y}\left[\mathbb{E}_X\left[\left(\mathbb{E}[y\mid x]-\mu(X)\right)^2\right]\right]=$$

$$=\mathbb{E}_{x,y}\left[\mathbb{E}_X\left[\left(\mathbb{E}[y\mid x]-\mathbb{E}_X\left[\mu(X)\right]+\mathbb{E}_X\left[\mu(X)\right]-\mu(X)\right)^2\right]\right]=$$

$$=\mathbb{E}_{x,y}\left[\mathbb{E}_X\left[\left(\mathbb{E}[y\mid x]-\mathbb{E}_X\left[\mu(X)\right]\right)^2\right]\right]+\mathbb{E}_{x,y}\left[\mathbb{E}_X\left[\left(\mathbb{E}_X\left[\mu(X)\right]-\mu(X)\right)^2\right]\right]+$$

$$+2\mathbb{E}_{x,y}\left[\mathbb{E}_X\left[\left(\mathbb{E}[y\mid x]-\mathbb{E}_X\left[\mu(X)\right]\right)\left(\mathbb{E}_X\left[\mu(X)\right]-\mu(X)\right)\right]\right].$$

Focus on this term

48

$$\mathbb{E}_X\left[\left(\mathbb{E}[y \mid x] - \mathbb{E}_X\left[\mu(X)\right]\right)\left(\mathbb{E}_X\left[\mu(X)\right] - \mu(X)\right)\right] =$$

$$\mathbb{E}_X \left[\left(\mathbb{E}[y \mid x] - \mathbb{E}_X \left[\mu(X) \right] \right) \left(\mathbb{E}_X \left[\mu(X) \right] - \mu(X) \right) \right] =$$

$$= \left(\mathbb{E}[y \mid x] - \mathbb{E}_X \left[\mu(X) \right] \right) \mathbb{E}_X \left[\mathbb{E}_X \left[\mu(X) \right] - \mu(X) \right] =$$

$$\mathbb{E}_X \Big[\big(\mathbb{E}[y \mid x] - \mathbb{E}_X [\mu(X)] \big) \big(\mathbb{E}_X [\mu(X)] - \mu(X) \big) \Big] =$$

$$= \big(\mathbb{E}[y \mid x] - \mathbb{E}_X [\mu(X)] \big) \mathbb{E}_X \Big[\mathbb{E}_X [\mu(X)] - \mu(X) \Big] =$$

$$= \big(\mathbb{E}[y \mid x] - \mathbb{E}_X [\mu(X)] \big) \Big[\mathbb{E}_X [\mu(X)] - \mathbb{E}_X [\mu(X)] \Big] =$$

$$\mathbb{E}_X\left[\left(\mathbb{E}[y\mid x] - \mathbb{E}_X\left[\mu(X)\right]\right)\left(\mathbb{E}_X\left[\mu(X)\right] - \mu(X)\right)\right] =$$

$$= \left(\mathbb{E}[y\mid x] - \mathbb{E}_X\left[\mu(X)\right]\right)\mathbb{E}_X\left[\mathbb{E}_X\left[\mu(X)\right] - \mu(X)\right] =$$

$$= \left(\mathbb{E}[y\mid x] - \mathbb{E}_X\left[\mu(X)\right]\right)\left[\mathbb{E}_X\left[\mu(X)\right] - \mathbb{E}_X\left[\mu(X)\right]\right] =$$

$$= 0.$$

$$\mathbb{E}_{x,y}\left[\mathbb{E}_X\left[\left(\mathbb{E}[y\mid x]-\mu(X)\right)^2\right]\right] =$$

$$= \mathbb{E}_{x,y}\left[\mathbb{E}_X\left[\left(\mathbb{E}[y\mid x]-\mathbb{E}_X\left[\mu(X)\right]+\mathbb{E}_X\left[\mu(X)\right]-\mu(X)\right)^2\right]\right] =$$

$$= \mathbb{E}_{x,y}\left[\mathbb{E}_X\left[\left(\mathbb{E}[y\mid x]-\mathbb{E}_X\left[\mu(X)\right]\right)^2\right]\right] + \mathbb{E}_{x,y}\left[\mathbb{E}_X\left[\left(\mathbb{E}_X\left[\mu(X)\right]-\mu(X)\right)^2\right]\right]+$$

$$+ 2\mathbb{E}_{x,y}\left[\mathbb{E}_X\left[\left(\mathbb{E}[y\mid x]-\mathbb{E}_X\left[\mu(X)\right]\right)\left(\mathbb{E}_X\left[\mu(X)\right]-\mu(X)\right)\right]\right].$$

0